


# Clock-Driven Scheduling

Ingo Sander

[ingo@kth.se](mailto:ingo@kth.se)

Liu: Chapter 5



IL2212 Embedded Software

1

## Outline

- Clock-driven scheduling
  - Notations and assumptions
  - Static, timer-driven cyclic schedules
  - Handling aperiodic jobs and sporadic jobs
  - Practical considerations
  - Pros and Cons



IL2212 Embedded Software

2



# Determinism

- Clock-driven scheduling requires a large amount of *determinism*
  - only few aperiodic and sporadic jobs
- Schedule can be calculated off-line



# Assumptions

1. There is a constant number  $n$  *periodic tasks* in the system
2. The parameters of all periodic tasks are known a priori
  - Variations in inter-release times of jobs are negligibly small
  - Each job in  $T_i$  is released  $p_i$  units of time after the previous job in  $T_i$
3. Each job  $J_{i,k}$  is ready for execution at its release time  $r_{i,k}$



# Assumptions

- There are aperiodic jobs released at unexpected time instants
  - aperiodic jobs are placed in special queue
  - new jobs are added to the queue without need to notify scheduler
  - when processor is available aperiodic jobs are scheduled
- There are no sporadic jobs (this assumption will be relaxed later)




## Static, Clock-Driven Scheduler

- Static schedule can be calculated off-line (all parameters are known at start)
  - complex algorithms can be used
  - amount of processor time allocated to each job is equal to its maximum execution time
  - static schedule guarantees that every job completes by its deadline as long as no job overruns



## Example

- Four independent periodic tasks:  $T_1 = (4, 1)$ ,  $T_2 = (5, 1.8)$ ,  $T_3 = (20, 1)$ ,  $T_4 = (20, 2)$
- Utilization =  $1/4 + 1.8/5 + 1/20 + 2/20 = 0.76$
- Hyperperiod = LCM (4, 5, 20, 20) = 20
- Possible schedule:

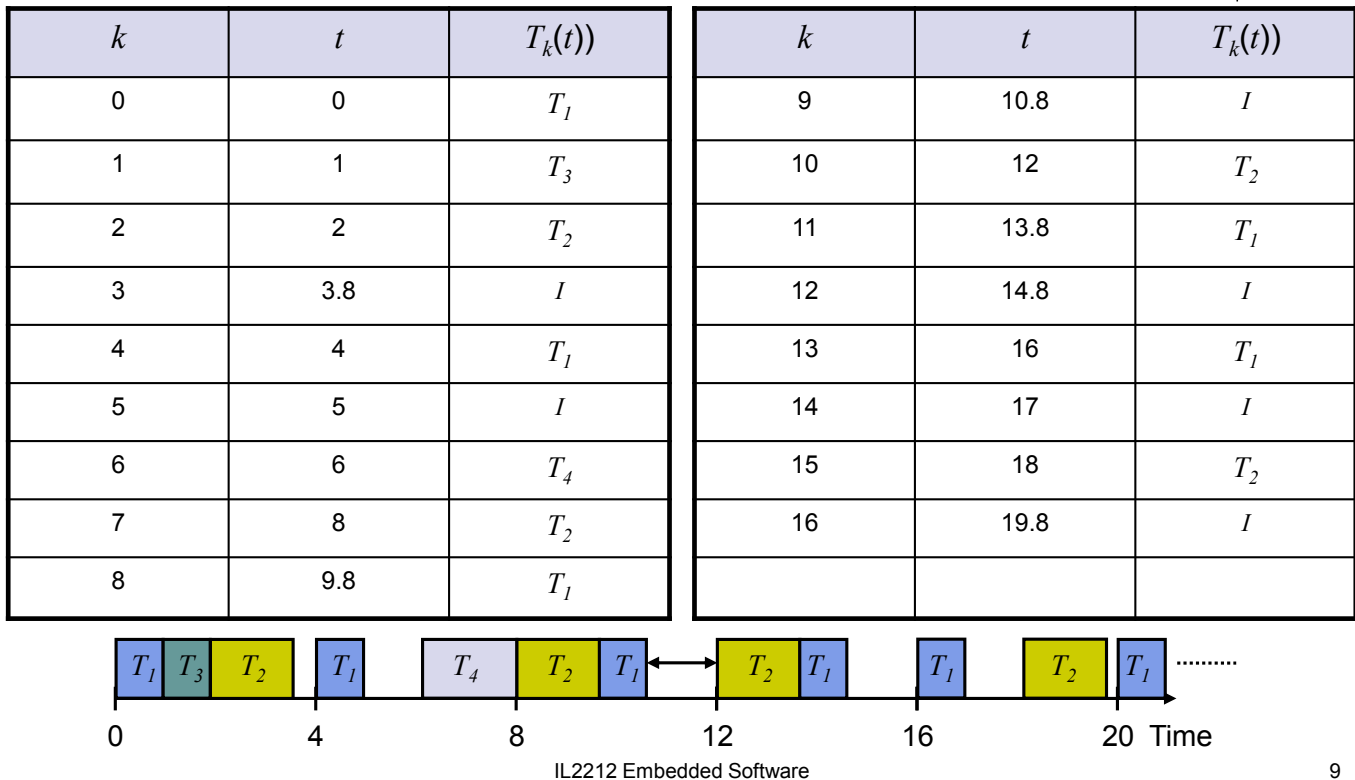


## Implementing a Cyclic Scheduler

- Store precomputed schedule as table
- Each entry  $(t, T_k(t))$  in table gives a *decision time*  $t_k$  at which a scheduling decision is made
- $T_k(t)$  can either be a task  $T_i$  or  $I$  (idle)
- Idle time can be used for aperiodic jobs



# Implementing a Cyclic Scheduler



9



# Implementing a Cyclic Scheduler

- Initialization
  - Tasks to be executed are created and sufficient memory is allocated
  - Code executed by the tasks is loaded into memory
- Scheduler is invoked on hardware timer interrupt
  - First interrupt at  $t_k = 0$
  - On receipt of an interrupt
    - Set next timer interrupt to  $t_{k+1}$
    - If  $T(t_k) = I$  and aperiodic job waiting start aperiodic job
    - otherwise schedule next job in task (and preempt aperiodic job, if there is one!)



# Structure of Cyclic Schedules

- Ad hoc table-driven schedules are flexible, but not efficient
  - relies on accurate timer interrupts and exact execution times of tasks
  - large scheduling overhead
  - intervals for aperiodic jobs are not spread out uniformly and maybe very short
  - interval timer needed



# Structure of Cyclic Schedules

- Frame-based approach
  - make scheduling decision periodically at certain intervals (*frames*)
  - execute a fixed number of jobs in every frame
  - each frame has a size of  $f$
  - preemption is only allowed at frame borders
  - the first job of every task is released at the beginning of a frame



# Structure of cyclic schedules

- Benefits
  - At the beginning of each frame
    - scheduler can check, if every job in frame has been released and is ready for execution
    - Scheduler can detect if there is any overrun or a missed deadline
  - Periodic timer instead of hardware timer can be used

For the theory how to calculate the optimal frame size check Liu, Section 5.3 (not part of the course)

IL2212 Embedded Software

13



## Cyclic Executive

- The term “cyclic executive” for a table-driven cyclic scheduler for all types of jobs in a multithreaded system
- Scheduling decisions are made at the beginning of each frame, triggered by timer interrupts



# Cyclic Executive

- During execution table entry for current frame is copied into current block
- Scheduler wakes up a job “periodic task server” that executes all job slices in the current block
- Then scheduler uses remaining time in frame for aperiodic jobs



# Scheduling Aperiodic Jobs


- So far aperiodic jobs have been scheduled in the background after all other job slices have been completed
  - Disadvantage: Average response time is long
- Average response time for aperiodic jobs can be improved by scheduling hard-real time jobs as late as possible without missing the deadline





# Slack Stealing

- Idea: Use slack to schedule aperiodic jobs before periodic jobs whenever possible!
- Implementation: Cyclic executive keeps track of slack and lets periodic task server execute aperiodic jobs as long as there is slack available




17



# Slack Stealing

- Interval timer is used
  - At beginning of frame timer is set to slack in frame
  - Whenever an aperiodic job executes slack is reduced
  - When timer expires, slack is consumed and aperiodic job is preempted



18



# Scheduling Sporadic Jobs

- Sporadic jobs have hard deadlines
- Minimum release times and maximum execution times are unknown a priori
  - No possibility to guarantee deadlines, when calculating schedule off-line before system start




# Scheduling Sporadic Jobs

- However the properties of a sporadic job are known, when the sporadic job is released
  - Acceptance test (at start of frame):
    - Sporadic job is only scheduled, if all scheduled jobs still meet their deadlines
    - Otherwise it is rejected
  - Accepted sporadic jobs can be scheduled using EDF



# Example Scheduling Sporadic Jobs




IL2212 Embedded Software

21



# Example Scheduling Sporadic Jobs

Off-line schedule




IL2212 Embedded Software

22




# Example Scheduling Sporadic Jobs



IL2212 Embedded Software

23

# Example Scheduling Sporadic Jobs



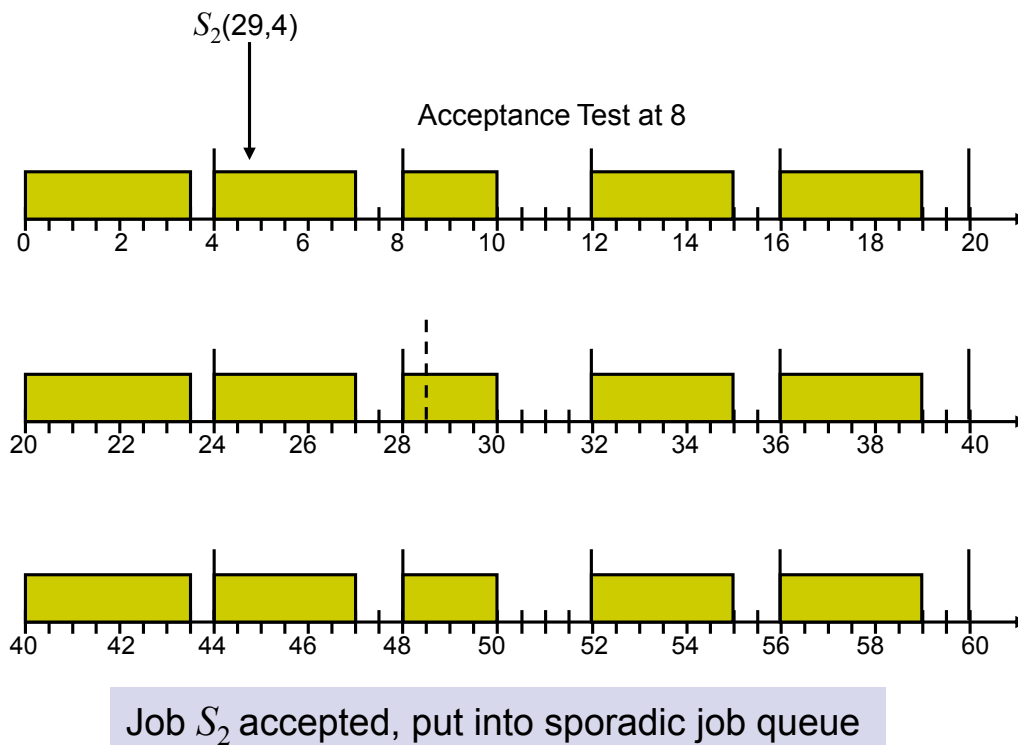
Job rejected: Not sufficient slack (Slack = 4)!

IL2212 Embedded Software

24



# Example Scheduling Sporadic Jobs

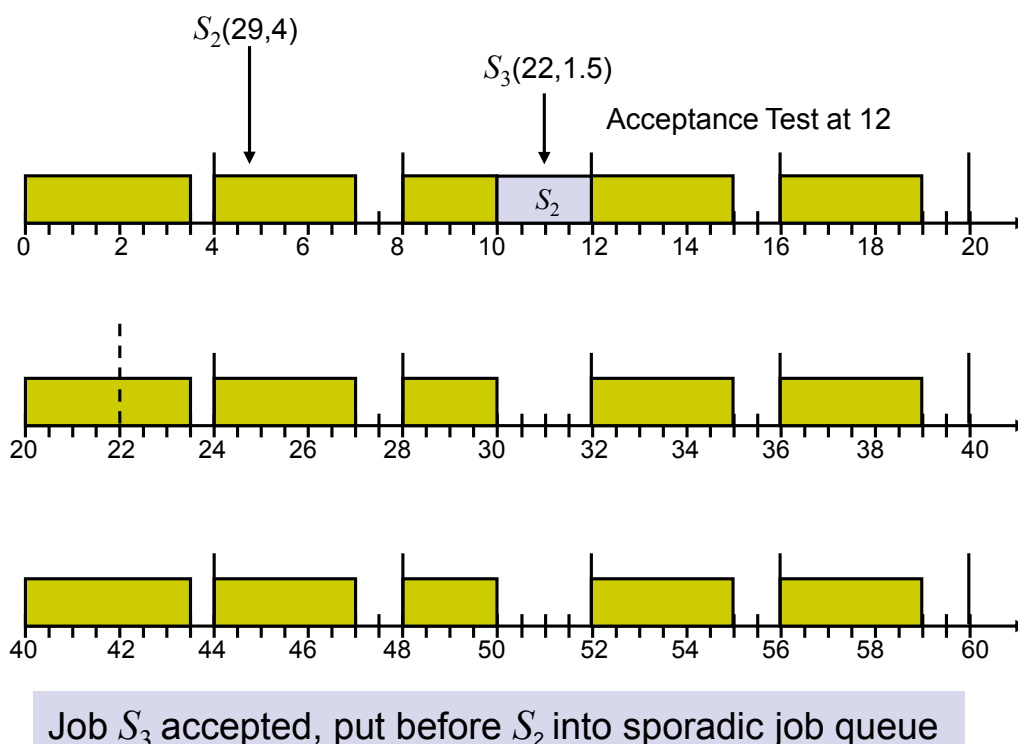


IL2212 Embedded Software

25



# Example Scheduling Sporadic Jobs

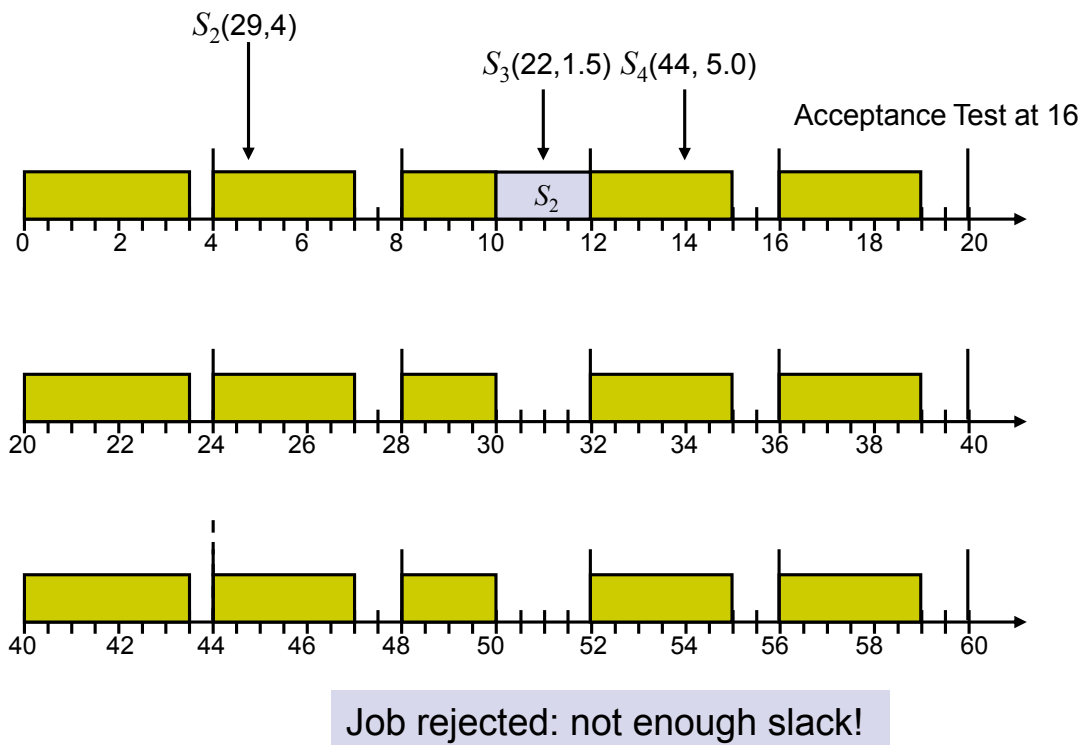


IL2212 Embedded Software

26



# Example Scheduling Sporadic Jobs




IL2212 Embedded Software

27



# Example Scheduling Sporadic Jobs



IL2212 Embedded Software

28



# Practical Considerations

- Handling Frame Overruns
  - Overruns may occur, when jobs execution time is longer than maximum execution time
  - Can be handled by
    - abort the overrun job and report the premature termination of the job
    - unfinished portion may execute as aperiodic job in a later frame
    - continue to execute the overrun job
      - but this may cause other jobs to be late, too!



# Practical Considerations

- Multiprocessor system
  - Constructing a feasible schedule for multiprocessors is more complex and time consuming than for uniprocessors
  - Since it is done off-line, exhaustive and complex heuristic algorithms can be used



# Advantages

- Conceptual simplicity
  - complex dependencies, communication delays can be considered when developing the schedule
  - no need for any synchronization mechanisms
  - schedule can be represented as tables that is used by the scheduler at run time
- Relatively easy to validate



# Disadvantages

- Inflexible, since schedule is computed off-line, small changes mean that new tables have to be generated
- Release times of jobs must be fixed
- A lot of information about jobs has to be known beforehand, so that the schedule can be pre-computed
- Difficult to get acceptable response times for aperiodic (soft real-time) jobs





# Summary

- Clock-driven schedulers schedule periodic tasks according to a cyclic schedule
- Task parameters must be known in advance
- Schedule can be calculated in advance
- Aperiodic and sporadic jobs can be scheduled, if they do not influence other scheduled jobs
- Applicable to static systems, with a small number of aperiodic jobs