

Priority-driven Scheduling of periodic tasks: EDF

OUTLINE

- Dynamic priority algorithms for scheduling periodic tasks
 - Optimality of EDF
 - Schedulability test

Ref: [Liu]

- Ch. 6 (pg. 124 – 129)

1

Schedulable utilization $U_{\text{EDF}} = 1$ if $D_i = p_i$

- **Theorem:** A system of independent preemptable periodic tasks with relative deadlines equal to their periods is feasible iff its total utilization U is less than or equal to 1.
- **Proof:**
 - **necessary condition (only-if):** if the task set is feasible, then $U \leq 1$;
this is the same as:
if $U > 1$, then the task set is not feasible,
which is obvious (we have already proved it);
 - **sufficient condition (if):** if $U \leq 1$, then the task set is feasible;
this is the same as:
if the task set is not feasible, then $U > 1$,
we show that if EDF fails to find a feasible schedule, then the total utilization must exceed 1.

2


$U_{EDF} = 1$ when $D_i = p_i$ (2)

- sufficient condition (if): if $U \leq 1$, then the task set is feasible;
which is the same as:
- if EDF fails to find a feasible schedule, then the total utilization must exceed 1.
 - assumptions:
 - at time t , job $J_{i,c}$ of task T_i (released at $r_{i,c}$) misses its deadline;
 - the processor never idles prior to t ;
 - current period of a task T_k is the period that starts before t and ends at or after t ; (current job: executed in the current period, with $r_{kc} < t$ and $d_{kc} \geq t$)
 - there are 2 cases to consider:
 - case 1: current period of every task begins at or after $r_{i,c}$;
 - case 2: current period of some task may start before $r_{i,c}$.

3

$U_{EDF} = 1$ when $D_i = p_i$ (3)

- case 1: current period of every task begins at or after $r_{i,c}$ (no current job with deadline after t is given any processor time before t)




- the number of full periods of task T_i up to time t is: $(t - \Phi_i)/p_i$
(this ratio is an integer number, since $(t - \Phi_i)$ is a multiple of p_i)
 - total processor time required by T_i up to time t is: $(t - \Phi_i)e_i/p_i$
- the number of full periods of any task $T_k \neq T_i$ up to time t is: $\lfloor (t - \Phi_k)/p_k \rfloor$
 - total processor time required by T_k up to time t is: $\lfloor (t - \Phi_k)/p_k \rfloor e_k$

4

$U_{EDF} = 1$ when $D_i = p_i$ (4)

• case 1(cont.):

- possible executions
- current periods



- total processor time required to complete all jobs (of T_i and all $T_{k \neq i}$) with deadline at or before t is: $(t - \phi_i)e_i/p_i + \sum_{k \neq i} \lfloor (t - \phi_k)/p_k \rfloor e_k$

- if $J_{i,c}$ misses its deadline at t , then the total processor time required by the task set before t exceeds the total available time t

$$\begin{aligned} t &< \frac{(t - \phi_i)e_i}{p_i} + \sum_{k \neq i} \lfloor \frac{t - \phi_k}{p_k} \rfloor e_k \\ &\leq t \cdot \frac{e_i}{p_i} + t \cdot \sum_{k \neq i} e_k/p_k \\ &= t \cdot U \\ &\Rightarrow U > 1 \end{aligned}$$


- this implies $U > 1$

5

$U_{EDF} = 1$ when $D_i = p_i$ (5)

- case 2: some current periods start before $r_{i,c}$

- possible executions
- current periods



T : Set of all tasks.

T' : tasks with $r'_c < r_{i,c}$ (current job starting before $r_{i,c}$) and with deadline $d'_c > t$ (T_i in the figure belongs to T')

$T^* = T - T'$: tasks with $r^*_c \geq r_{i,c}$ (T_i belongs to T^*), or with $r^*_c < r_{i,c}$ and $d'_c = t$.

$T = T' + T^*$

- t_1 : last point in time before t when some current job in T' is executing ($t_1 \leq r_{i,c}$).
- current jobs in T' may have executed before t_1 ; but no job in T' is given any processor time in $[t_1, t]$.
- the only jobs executed in $[t_1, t]$ belong to T^* , e.g. $J_{1,f}$ and $J_{n,g}$ in the figure (not in the current period), or jobs with current period ending at t (e.g. $J_{i,c}$).

6

$U_{EDF} = 1$ when $D_i = p_i$ (6)

• case 2 (cont):

- let ϕ_k' be the release time of the first job of task T_k in T^* starting from t_{-1}

→ IN THE SEGMENT $[t_{-1}, t]$:

- processor time required by $J_{i,c}$ and other jobs of T_i is:

$$(t - t_{-1} - \phi_i')e_i/p_i$$

(= me_i with m integer ≥ 1);


- none of the current jobs with deadline after t is given any processor time;

- the only jobs executed are: jobs of T_i released at or after t_{-1} (including $J_{i,c}$) and other jobs (of $T'' = T^* - T_i$) with deadline before t (not in the current period) or at t ;

- total processor time required by all these other jobs (of tasks in T'') is:

$$\sum_{k \in T''} \lfloor (t - t_{-1} - \phi_k')/p_k \rfloor e_k$$

7



$U_{EDF} = 1$ when $D_i = p_i$ (7)

• case 2 (cont):

■ ■ possible
■ ■ executions
--- current periods

- if $J_{i,c}$ misses its deadline at t , then total processor time required by $J_{i,c}$ and all other jobs (of T'') in $[t_{-1}, t]$ exceeds the total available time $t - t_{-1}$:

$$t - t_{-1} < (t - t_{-1} - \phi_i')e_i/p_i + \sum_{k \in T''} \lfloor (t - t_{-1} - \phi_k')/p_k \rfloor e_k$$

$$t - t_{-1} < (t - t_{-1})e_i/p_i + (t - t_{-1})\sum_{k \in T''} e_k/p_k$$


$$t - t_{-1} < (t - t_{-1})(e_i/p_i + \sum_{k \in T''} e_k/p_k + \sum_{j \in T'} e_j/p_j)$$

$$t - t_{-1} < (t - t_{-1})U$$

$$T = T^* + T' = T_i + T'' + T'$$


$$U > 1$$

8



$$U_{\text{EDF}} = 1 \text{ when } D_i = p_i \text{ (8)}$$

- What about the assumption that the processor never idles?



9

$$\text{EDF: } D_i \geq p_i$$

- From the above theorem, it follows that:
 1. A system of independent preemptable periodic tasks with **relative deadlines longer than their periods** is feasible iff their total utilization is less than or equal to 1.
 2. The schedulable utilization U_{EDF} of the EDF algorithm for independent preemptable periodic tasks with **relative deadlines equal to or longer than their periods** is 1.


10

EDF: $D_i < p_i$ (for some tasks)

- When the **relative deadlines** of some tasks are **less than their period**, the system may not be feasible, even when its total utilization is less than 1.

- ex.: $T_1=(2,0.9)$; $T_2=(5,2.3,3)$ $U = 0.9/2+2.3/5 = 0.91 < 1$

- nevertheless the system is not feasible:



11

EDF: DENSITY

- To take into account the possibility that $D_i < p_i$ for some tasks, we define the **DENSITY** of a periodic task T_i as the ratio between its execution time e_i and the minimum between its period p_i and its relative deadline D_i :

$$\delta_i = e_i / \min\{p_i, D_i\}$$

- the density Δ of a system of n periodic tasks is the sum of the densities of all its tasks:

$$\Delta = \sum_{i=1 \text{ to } n} e_i / \min\{p_i, D_i\}$$

- since $e_i / \min\{p_i, D_i\} \geq e_i / p_i$ for any i , the density of a system of n periodic tasks is never less than its total utilization:

$$\Delta \geq U$$

- if the density of a system is >1 the system may not be feasible

- ex.: $T_1=(2,0.9)$; $T_2=(5,2.3,3)$; $\Delta = 0.9/2+2.3/3 = 7.3/6 > 1$

12

EDF: $D_i \leq p_i$

- When $D_i \leq p_i$ for all tasks of a task set, its schedulability may be verified with the following:
- **Theorem:** a system T of independent, preemptable, periodic tasks, with relative deadlines not longer than their periods, can be feasibly scheduled in one processor if its total density is less than or equal to 1.
- **Proof:**
 - Let's consider the set $T = \{T_1, T_2, \dots, T_n\}$ with $D_k \leq p_k$ for any k .
 - the density of a task T_k is: $\delta_k = e_k / \min\{D_k, p_k\}$
 - the total density of the task set T is: $\Delta = \sum_{k=1 \text{ to } n} \delta_k$
 - we will prove that if the task set T is not schedulable, then $\Delta > 1$.

13

EDF: $D_i \leq p_i$ (2)

- let t be the time instant in which a job $J_{i,c}$ of task T_i , released at $r_{i,c}$, misses its deadline: $t = r_{i,c} + D_i$ and let's put the time origin $t = 0$ at the last instant before t in which the processor was idle or executing jobs with deadline after t (we may ignore what happened before $t = 0$).
- then in $[0, t]$ the processor is always executing jobs released at or after 0 and having deadlines $\leq t$; if $J_{i,c}$ misses its deadline at t , the total time requested by these jobs is greater than t :

$$t < TR_i + TR_o$$

where:


TR_i = total time requested by all jobs of T_i released before t

TR_o = total time requested by all other jobs of T with deadline $\leq t$

- let's compute these two amounts of time:

14

EDF: $D_i \leq p_i$ (3)



- let's compute the two amounts of time TR_i and TR_o :
 - the first job of T_i is released at Φ_i ; the last job of T_i before t is released at $r_{i,c} = t - D_i$; the number of full periods of T_i in the time interval $[\Phi_i, r_{i,c}]$ is: $(r_{i,c} - \Phi_i)/p_i$ (this ratio is an integer number); the number of jobs of T_i released in $[\Phi_i, r_{i,c}]$ is:
 $(r_{i,c} - \Phi_i)/p_i + 1 = (t - D_i - \Phi_i)/p_i + 1$
 - $TR_i = [(t - D_i - \Phi_i)/p_i + 1]e_i$
 - the first job of $T_k \in T_o$ is released at Φ_k ; the last job of T_k released (at $r_{k,c}$) before t is executed before t only if its deadline $d_{k,c}$ is not later than t : $d_{k,c} = r_{k,c} + D_k \leq t$ i.e. only if $r_{k,c} \leq t - D_k$
 - then the jobs of T_k executed before t are those released in the time interval $[\Phi_k, t - D_k]$; their number is: $\lfloor (t - D_k - \Phi_k)/p_k \rfloor + 1$
 - $TR_o = \sum_{k \in o, k \neq i} [\lfloor (t - D_k - \Phi_k)/p_k \rfloor + 1]e_k$

15

EDF: $D_i \leq p_i$ (4)

- if $J_{i,c}$ misses its deadline at t then the processor time requested by all jobs of T up to the time t is greater than t :

$$t < [(t - D_i - \Phi_i)/p_i + 1]e_i + \sum_{k \in o, k \neq i} [\lfloor (t - D_k - \Phi_k)/p_k \rfloor + 1]e_k$$

$$\leq [(t - D_i)/p_i + 1]e_i + \sum_{k \in o, k \neq i} [(t - D_k)/p_k + 1]e_k$$
 - consider the first term: $[(t - D_i)/p_i + 1]e_i$
 - if $D_i = p_i$, the term reduces to: $te_i/p_i = t\bar{\delta}_i$
 - if $D_i < p_i$, the term is $< [(t - D_i)/D_i + 1]e_i = te_i/D_i = t\bar{\delta}_i$
 - in any case the first term is $\leq t\bar{\delta}_i$: $[(t - D_i)/p_i + 1]e_i \leq t\bar{\delta}_i$
 - similarly for the second term:
 $\sum_{k \in o, k \neq i} [(t - D_k)/p_k + 1]e_k \leq \sum_{k \in o, k \neq i} t\bar{\delta}_k = t\sum_{k \in o, k \neq i} \bar{\delta}_k$
- then the inequality:

$$t < [(t - D_i - \Phi_i)/p_i + 1]e_i + \sum_{k \in o, k \neq i} [\lfloor (t - D_k - \Phi_k)/p_k \rfloor + 1]e_k$$
 becomes: $t < t\bar{\delta}_i + t\sum_{k \in o, k \neq i} \bar{\delta}_k = t\sum_{k=1 \text{ to } n} \bar{\delta}_k = t\Delta$
- $t < t\Delta$ implies $\Delta > 1$.
- the theorem gives only a sufficient condition for the schedulability of task sets with $D_i \leq p_i$: a task set may be schedulable even when $\Delta > 1$


16

Feasible system with $\Delta > 1$

- Example: $T_1 = (2, 0.6, 1)$; $T_2 = (5, 2.3)$

- Density $\Delta = 0.6/1 + 2.3/5 = 0.6 + 0.46 = 1.06 > 1$

- EDF schedule:



- The system is feasible with $\Delta > 1$

17

EDF: arbitrary deadlines

- From the previous two theorems, follows the following general:
- **Theorem:** a system T of independent, preemptable, periodic tasks, with arbitrary relative deadlines and periods, can be feasibly scheduled in one processor if its density is less than or equal to 1 (if $\Delta \leq 1$).
- **Proof:**
 - case 1. $D_k \geq p_k$ for all k :
in this case $\delta_k = e_k / \min\{D_k, p_k\} = e_k / p_k = u_k$ for all k ; then $U = \Delta$; $\Delta \leq 1$ implies $U \leq 1$, which, according to the first of the last two theorems, is a necessary and sufficient condition for schedulability;
 - case 2. $D_k < p_k$ for some k :
according to the last theorem, $\Delta \leq 1$ is a sufficient condition for schedulability.

18

Schedulability test

SCHEDULABILITY TEST:

- For a system $T = \{T_1, T_2, \dots, T_n\}$ of independent periodic tasks, we are given:
 - p_i, e_i, D_i (period, execution time and relative deadline) of every T_i
 - a priority driven algorithm to schedule T preemptively on one processor
- Determine whether all the deadlines of every task are always met
- If the schedulability test is efficient, it can be used as an **on-line acceptance test**.

19

Schedulability test for EDF

- When the scheduling algorithm is EDF, the 2 theorems presented previously (about optimality and density) provide the theoretical basis for a simple general schedulability test:
- Check the inequality:

$$\sum_{k=1 \text{ to } n} e_k / \min(D_k, p_k) \leq 1$$

- if it is satisfied, the system is schedulable
- if it is not satisfied then:
 - if $D_k \geq p_k$ for all k , then the inequality reduces to **$U \leq 1$** which is both a necessary and sufficient condition; if the inequality is not satisfied \Rightarrow the system is not schedulable;
 - if $D_k < p_k$ for some k , then the inequality is only a sufficient condition, \Rightarrow the system may not be schedulable.

20

Robustness of schedulability test

- Robustness against variations of execution times:
 - a system which passes the schedulability test remains feasible even if some jobs have a **shorter execution time** (recall that the execution of independent preemptable jobs on one processor is predictable): no anomalous behaviour may take place.
- Robustness against variations of interrelease times:
 - the theorem in slide 2 has been proved taking into account the total demand of processor time by the jobs of each task, regardless of the actual values of their periods: if the interrelease time of some jobs is longer, then the total demand of processor time becomes smaller; as a consequence:
 - a system which passes the schedulability test remains feasible even if some jobs have a **longer interrelease time**.

21

Schedulability test as a design aid

- The simple schedulability test:

$$\sum_{k=1 \text{ to } n} e_k / \min(D_k, p_k) \leq 1$$

may be used **during system design** to guide the choices of periods of tasks related to their execution times:

- adding a new task to the systems without causing the other tasks to miss any deadline may require to adjust task periods in such a way that the test is satisfied;
 - if the other task periods cannot be modified, the test inequality provides the minimum value for the period of the new task that guarantees feasibility of the system.
- The same simple schedulability test:

$$\sum_{k=1 \text{ to } n} e_k / \min(D_k, p_k) \leq 1$$

may be used as an **online acceptance test** for adding new tasks.

22

EDF: time demand analysis

- The simple sufficient condition for EDF schedulability:

$$\Delta \leq 1$$

may be conveniently used during the design phase to validate a set of periodic tasks with arbitrary deadlines.

- When $D_k \leq p_k$ for any $k = 1, \dots, n$, a **more accurate schedulability test** may be performed, based on the evaluation of the following 2 functions:
 - **$W(t)$: time demand**: processor time demanded by all jobs released before t , i.e. in the time interval $[0, t)$;
 - **$V(t)$: demanded time**: processor time that must have been provided to meet the demand of all jobs having deadlines before or at t , i.e. in the time interval $[0, t]$.

23

Processor time demand functions

- the number of jobs of a task T_k released in the time interval $[0, t)$ is $\lceil (t - \Phi_k)/p_k \rceil$; the total processor time demanded by all these jobs is $\lceil (t - \Phi_k)/p_k \rceil e_k$; then, for the whole task set:

$$W(t) = \sum_{k=1 \text{ to } n} \lceil (t - \Phi_k)/p_k \rceil e_k$$

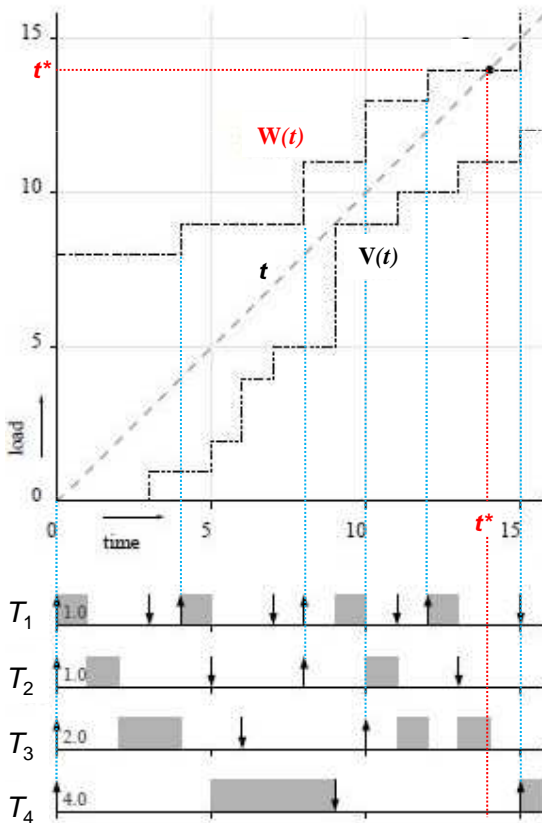
- to evaluate $V(t)$:
 - the number of full periods of T_k in $[0, t]$ is $\lfloor (t - \Phi_k)/p_k \rfloor$; the number of jobs of T_k released in $[0, t]$ is $\lfloor (t - \Phi_k)/p_k \rfloor + 1$;
 - since $D_k \leq p_k$, all of these jobs except the last one (let's call $r_{k,c}$ its release time) must certainly have been completed before t , while the last one must have been executed before t only if its deadline $d_{k,c} = r_{k,c} + D_k$ is not later than t , i.e. if $r_{k,c} + D_k \leq t$, or $r_{k,c} \leq t - D_k$
 - then the jobs of T_k that must have been completed before t are those released in $[\Phi_k, t - D_k]$; their number is $\lfloor (t - D_k - \Phi_k)/p_k \rfloor + 1$;
 - the processor time that must have been used by all these jobs is

$$V(t) = \sum_{k=1 \text{ to } n} (\lfloor (t - D_k - \Phi_k)/p_k \rfloor + 1) e_k$$

24

Example: time demand $W(t)$

ex: EDF schedule of $T_1 = (4,1,3)$; $T_2 = (8,1,5)$; $T_3 = (10,2,6)$; $T_4 = (15,4,9)$



$W(t)$ is a step function that increases of a quantity e_k whenever a job of T_k is released: i.e. at $t = ip_k$, for $k = 1, 2, 3, 4$ and $i = 0, 1, 2, 3, \dots$

The 45° line represents the available processor time.


The difference between $W(t)$ and this line represents the amount of processor time necessary to complete all jobs that have been released up to t . While all of these jobs have not completed, $W(t)$ lies above the 45° line.

When $W(t)$ touches the 45° line, $t^* = W(t^*)$: the available processor time has been sufficient to satisfy the demand of all jobs released up to t^* (the processor starts an idle period, unless a new job is released at t^*).

25

Example: demanded time $V(t)$

ex: EDF schedule of $T_1 = (4,1,3)$; $T_2 = (8,1,5)$; $T_3 = (10,2,6)$; $T_4 = (15,4,9)$



$V(t)$ is a step function that increases of a quantity e_k at each instant in which a job of T_k has its deadline: i.e. at $t = id_k$, for $k = 1, 2, 3, 4$ and $i = 1, 2, 3, \dots$

As long as $V(t)$ lies below the 45° line, $V(t) \leq t$, then all deadlines up to t have been met.

If $V(t)$ crosses the 45° line at $t = t^\wedge$, $V(t^\wedge) > t^\wedge$, the available processor time has not been sufficient to satisfy the demand of all jobs with deadlines before or at t^\wedge and the deadline at t^\wedge is missed.

(In the figure no deadline is missed.)

26

Demanded time function $V(t)$

$$V(t) = \sum_{k=1}^n (\lfloor (t - D_k - \Phi_k)/p_k \rfloor + 1) e_k$$

- A set of tasks with $D_k \leq p_k$ (for any k) is feasible if and only if the time demanded by its jobs never exceeds the available time, i.e. iff $V(t) \leq t$ (for any t)
- In case all tasks have zero phase ($\Phi_k = 0$, for any k), the demanded time function becomes:

$$V(t) = \sum_{k=1}^n (\lfloor (t - D_k)/p_k \rfloor + 1) e_k = \sum_{k=1}^n \lfloor (t - D_k + p_k)/p_k \rfloor e_k$$

- If $V(t) \leq t$ throughout the first hyperperiod, then the time demand of all jobs released in the hyperperiod never exceeds the available processor time; since the schedule repeats itself identical in each hyperperiod, the task set is feasible.
- Since $V(t)$ increases at each deadline d_k and remains constant until the next deadline d_{k+1} , the condition $V(t) \leq t$ may be verified only for values of t equal to jobs' deadlines in the first hyperperiod:

$$\sum_{k=1}^n \lfloor (t - D_k + p_k)/p_k \rfloor e_k \leq t \text{ for } t = d_k (d_k \leq H)$$

27

Demanded time $V(t)$ and utilization U

- For tasks sets with $D_k = p_k$ (for all k), the condition $V(t) \leq t$ becomes:

$$V(t) = \sum_{k=1}^n \lfloor (t - D_k + p_k)/p_k \rfloor e_k = \sum_{k=1}^n \lfloor t/p_k \rfloor e_k \leq t$$

which is equivalent to the condition $U \leq 1$.

- Proof:

- part 1: if $U \leq 1$, then $V(t) \leq t$ for all $t > 0$.

if $U \leq 1$, then, for all t , we have:

$$t \geq Ut = \sum_{k=1}^n (t/p_k) e_k \geq \sum_{k=1}^n \lfloor t/p_k \rfloor e_k = V(t)$$

- part 2 if $V(t) \leq t$ for all $t > 0$, then $U \leq 1$.


we show that if $U > 1$, then $V(t) > t$ for at least one value of t .

if $U > 1$, then for $t^* = \text{lcm}(p_1, \dots, p_n)$ we have:

$$t^* < Ut^* = \sum_{k=1}^n (t^*/p_k) e_k = \sum_{k=1}^n \lfloor t^*/p_k \rfloor e_k = V(t^*)$$

28

Example: $W(t)$ and $V(t)$



$$T_1 = (2, 0.6, 1); T_2 = (5, 2.3)$$

Density $\Delta = 1.06 > 1$ - EDF schedule

The time demand function $W(t)$ crosses the 45° line at $t = 3.5$ and the system becomes idle; when a new job ($J_{1,3}$) is released at $t = 4$, its time demand (of $e_1 = 0.6$ time units) starts at $t = 4$; thus the new time demand steps up to $4 + 0.6$: $W(4) = 4.6$

The same happens each time a new job is released after the system has become idle: at $t = 5$, $t = 8$, $t = 10$.

The demanded time function $V(t)$ lies below the 45° line during the whole hyperperiod of the task set: therefore the task set is feasible.